

# Module 11: External memory

## CS 240 - Data Structures and Data Management

Mark Petrick

Based on lecture notes by many previous cs240 instructors

David R. Cheriton School of Computer Science, University of Waterloo

Fall 2017

# Different levels of memory

Current architectures:

- registers (very fast, very small)
- cache L1, L2 (still fast, less small)
- main memory
- external memory: disk (slow, very large)

General question: how to adapt our algorithms to take the memory hierarchy into account, avoiding transfers as much as possible?

# Review: dictionary ADT

A *dictionary* is a collection of *items*, each of which contains

- a *key*
- some *data*,

and is called a *key-value pair* (KVP). Keys can be compared and are (typically) unique.

Operations:

- *search*(*k*)
- *insert*(*k*, *v*)
- *delete*(*k*)
- optional: *join*, *isEmpty*, *size*, etc.

# Dictionaries in external memory

Tree-based data structures have poor *memory locality*:  
If an operation accesses  $m$  nodes, then it must access  $m$  spaced-out memory locations.

**Observation:** Accessing a single location in *external memory* (e.g. hard disk) automatically loads a whole block (or “page”).

- In an AVL tree,  $\Theta(\log n)$  pages are loaded in the worst case.
- Better solution: B-trees

## 2-3 Trees

A 2-3 Tree is like a BST with additional structural properties:

- Every internal node either contains *one KVP* and *two children*, or *two KVPs* and *three children*.
- The leaves are *NIL* (do not store keys)
- All the leaves are at the same level.

Searching through a 1-node is just like in a BST.

For a 2-node, we must examine both keys and follow the appropriate path.

## Insertion in a 2-3 tree

First, we search to find the lowest internal node where the new key belongs.

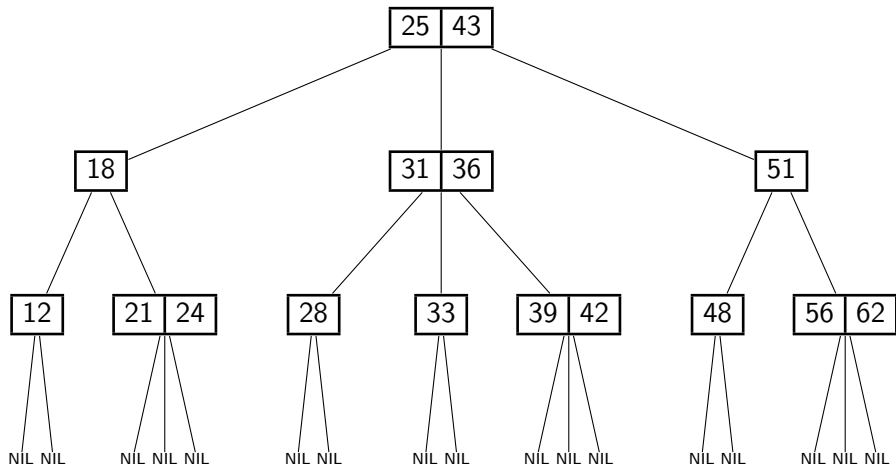
If the node has only 1 KVP, just add the new one to make a 2-node.

Otherwise, order the three keys as  $a < b < c$ .

Split the node into two 1-nodes, containing  $a$  and  $c$ ,  
and (recursively) insert  $b$  into the parent along with the new link.

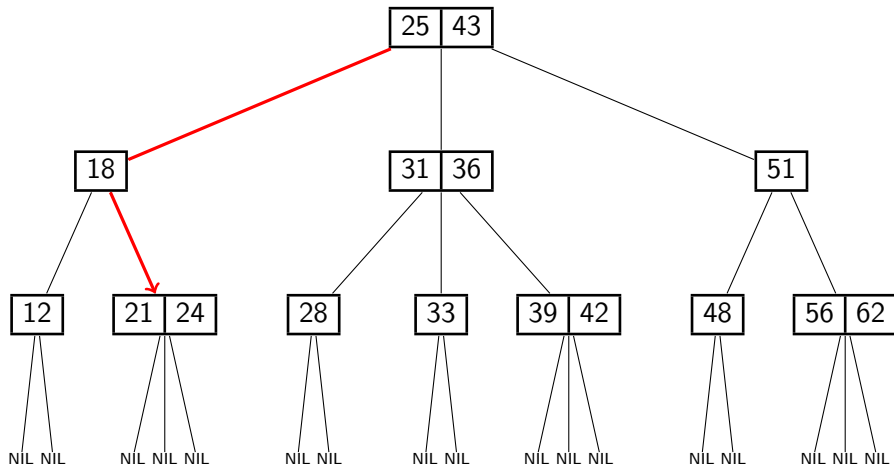
## 2-3 Tree Insertion

**Example:** *insert*(19)



## 2-3 Tree Insertion

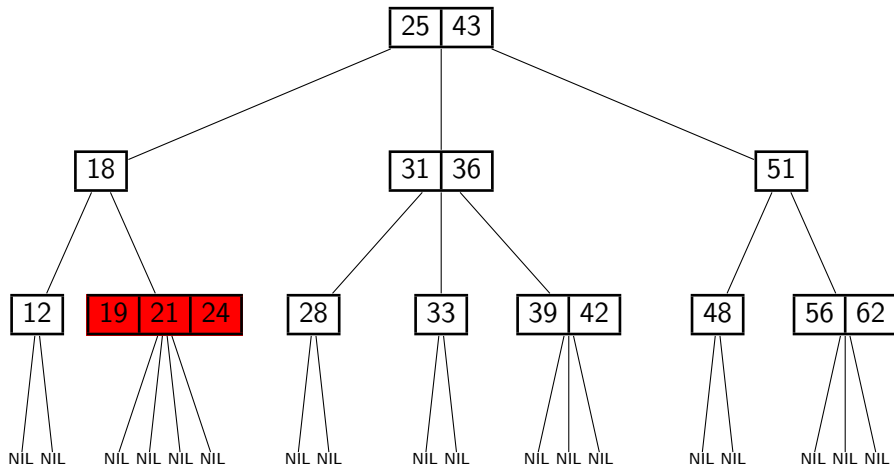
**Example:** *insert*(19)





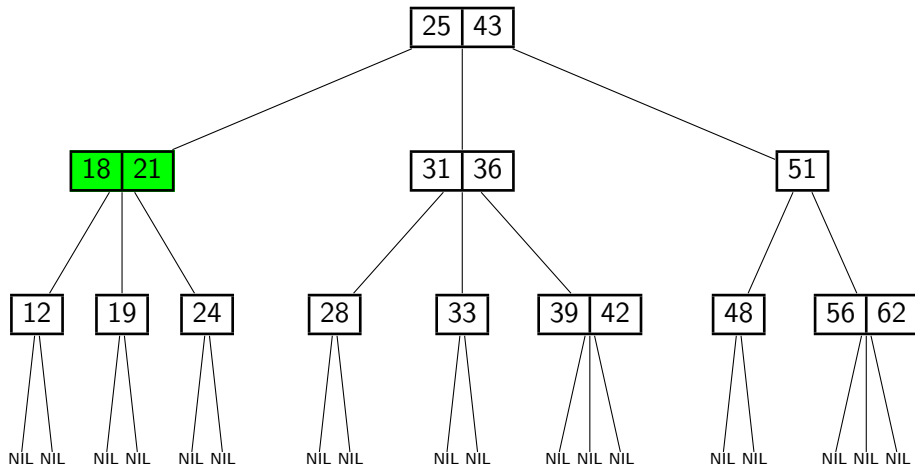
## 2-3 Tree Insertion

**Example:** *insert*(19)



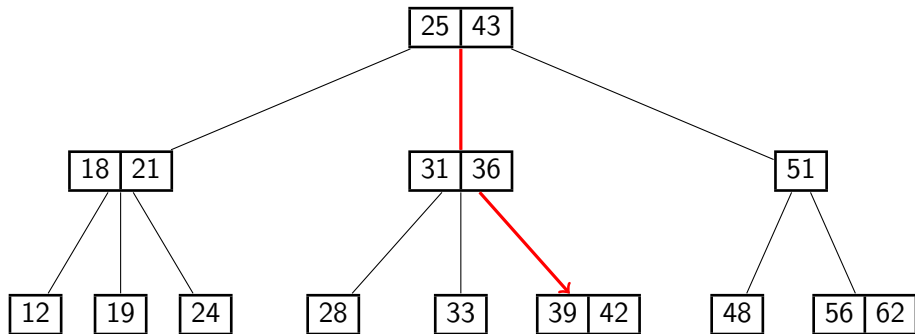
## 2-3 Tree Insertion

**Example:** *insert*(19)



## 2-3 Tree Insertion

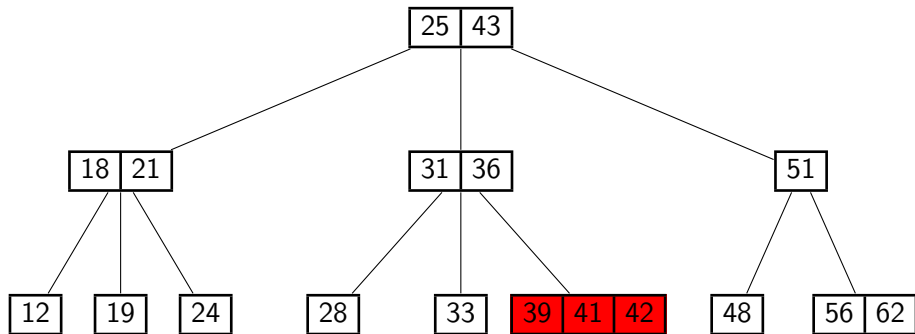
**Example:** *insert*(41)



(NIL-leaves not shown to simplify picture)

## 2-3 Tree Insertion

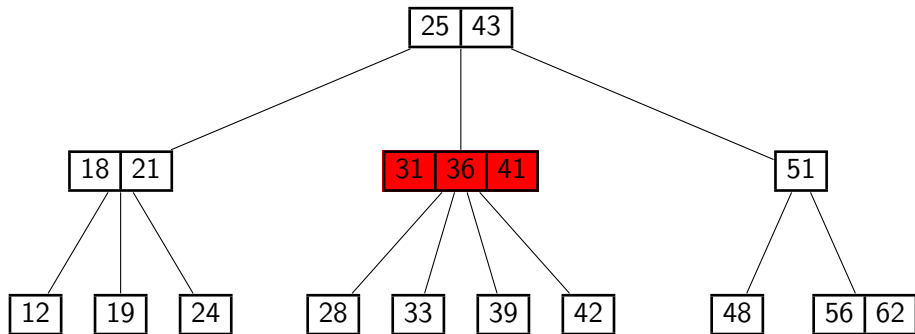
**Example:** *insert*(41)



(NIL-leaves not shown to simplify picture)

## 2-3 Tree Insertion

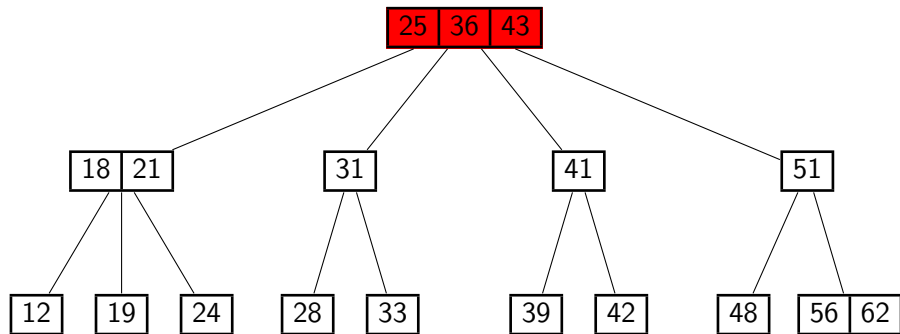
**Example:** *insert*(41)



(NIL-leaves not shown to simplify picture)

## 2-3 Tree Insertion

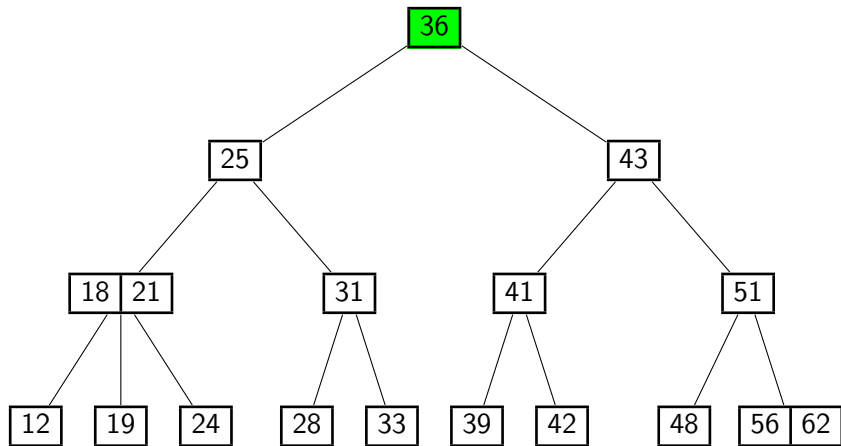
**Example:** *insert*(41)



(NIL-leaves not shown to simplify picture)

## 2-3 Tree Insertion

**Example:** *insert*(41)



(NIL-leaves not shown to simplify picture)

## Deletion from a 2-3 Tree

As with BSTs and AVL trees, we **first swap** the KVP with its successor, so that we always delete from a leaf.

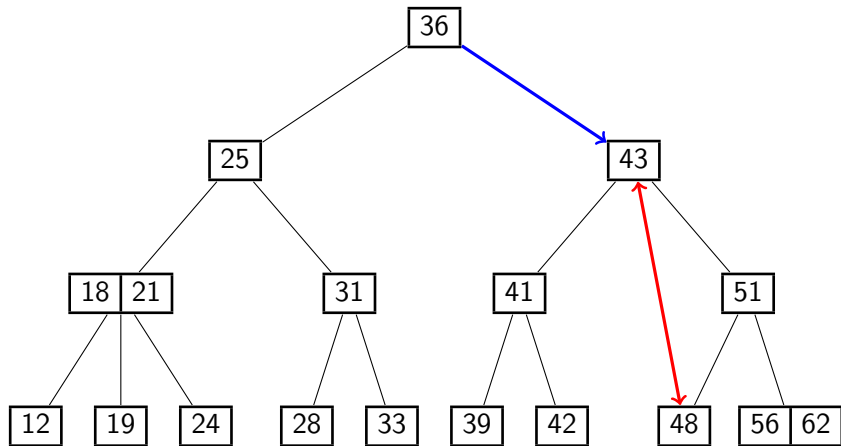
Say we're deleting KVP  $x$  from a node  $V$ :

- If  $V$  is a 2-node, just delete  $x$ .
- Else if  $V$  has a 2-node *immediate* sibling  $U$ , perform a *transfer*:  
Put the “intermediate” KVP in the parent between  $V$  and  $U$  into  $V$ , and replace it with the adjacent KVP from  $U$ .
- Otherwise, we *merge*  $V$  and a 1-node sibling  $U$ :  
Remove  $V$  and (recursively) delete the “intermediate” KVP from the parent, adding it to  $U$ .



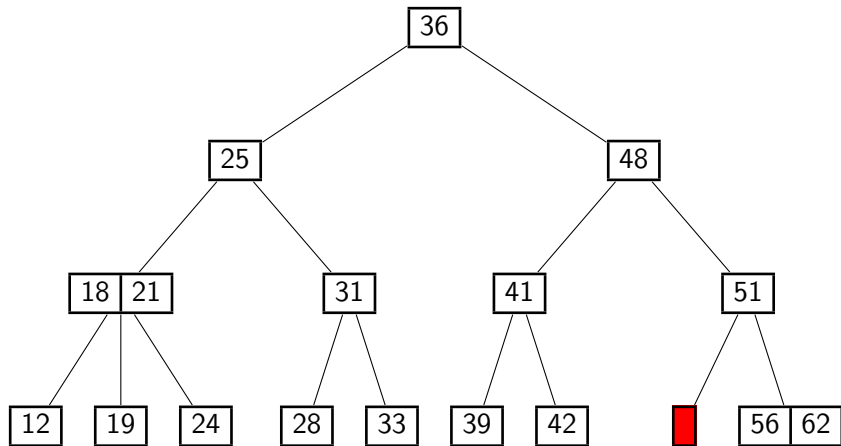
## 2-3 Tree Deletion

**Example:** *delete*(43)




## 2-3 Tree Deletion

**Example:** *delete*(43)




## 2-3 Tree Deletion

**Example:** *delete*(43)




## 2-3 Tree Deletion

**Example:** *delete*(19)



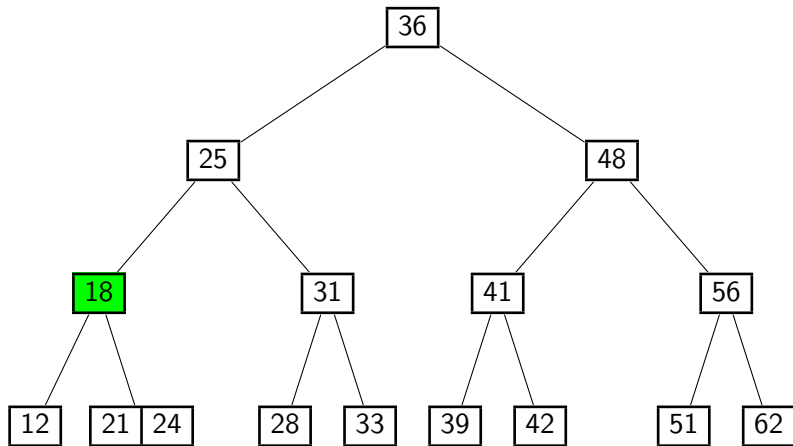
## 2-3 Tree Deletion

**Example:** *delete*(19)



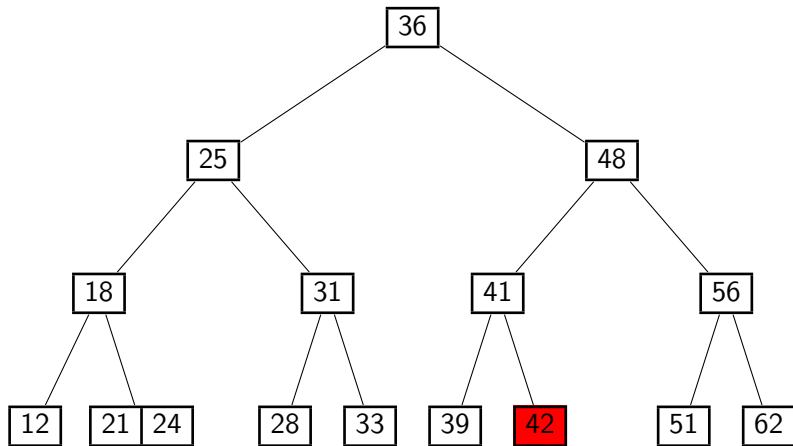
## 2-3 Tree Deletion

**Example:** *delete*(19)



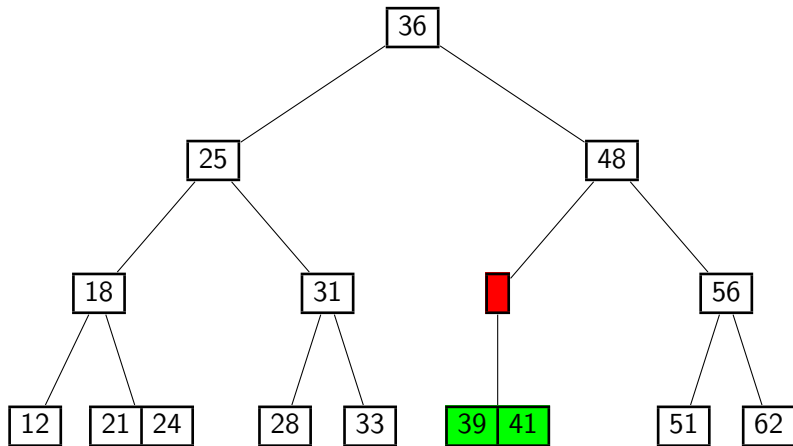
## 2-3 Tree Deletion

**Example:** *delete*(42)



## 2-3 Tree Deletion

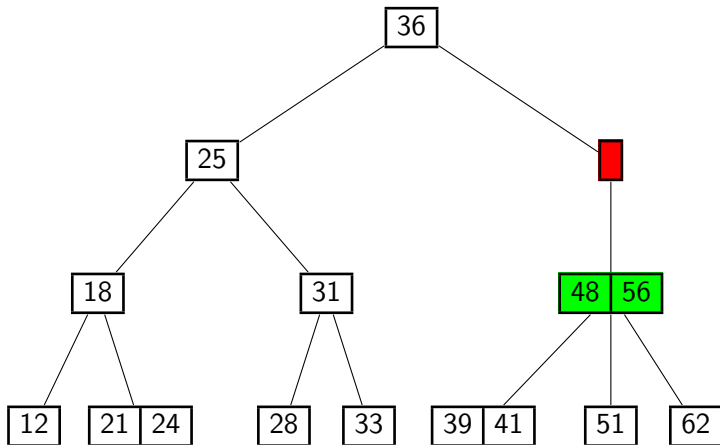
**Example:** *delete*(42)





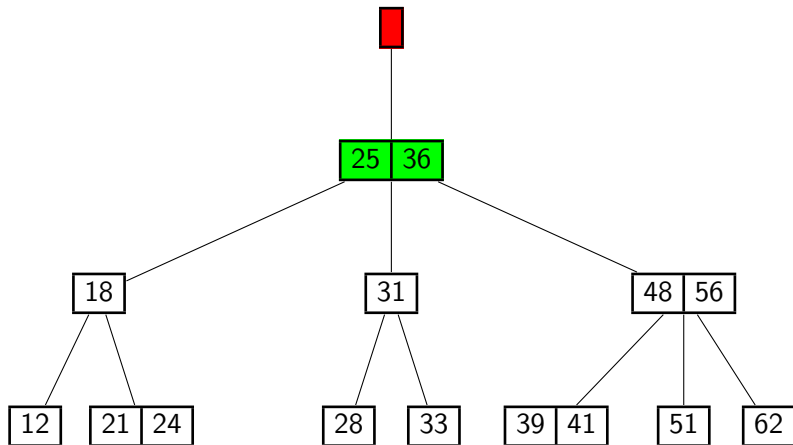
## 2-3 Tree Deletion

**Example:** *delete*(42)



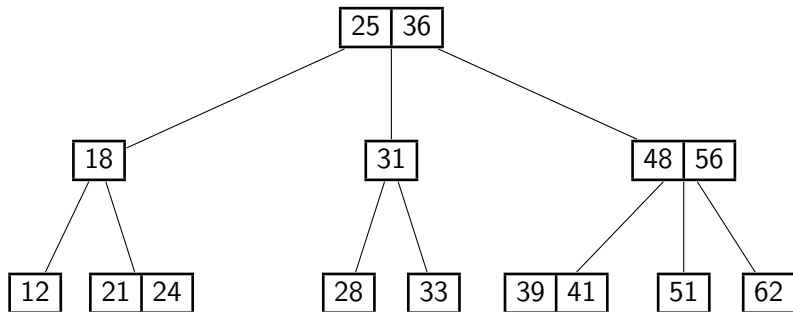
## 2-3 Tree Deletion

**Example:** *delete*(42)



## 2-3 Tree Deletion

**Example:** *delete*(42)



# B-Trees

The 2-3 Tree is a specific type of  $(a, b)$ -tree:

An  $(a, b)$ -tree of order  $M$  is a search tree satisfying:

- Each internal node has at least  $a$  children, unless it is the root.  
The root has at least 2 children.
- Each internal node has at most  $b$  children.
- If a node has  $k$  children, then it stores  $k - 1$  key-value pairs (KVPs).
- Leaves store no keys and are at the same level. for 3, except root node

A  $B$ -tree of order  $M$  is a  $(\lceil M/2 \rceil, M)$ -tree.

A 2-3 tree has  $M = 3$ .

*search*, *insert*, *delete* work just like for 2-3 trees.

## Height of a B-tree

What is the least number of KVPs in a height- $h$  B-tree?

(Height = # levels **not** counting the dummy-level  $-1$ )

Level	Nodes	Links/node	KVP/node	KVPs on level
0	1	2	1	1
1	2	$M/2$	$M/2 - 1$	$2(M/2 - 1)$
2	$2(M/2)$	$M/2$	$M/2 - 1$	$2(M/2)(M/2 - 1)$
3	$2(M/2)^2$	$M/2$	$M/2 - 1$	$2(M/2)^2(M/2 - 1)$
...	...	...	...	...
$h$	$2(M/2)^{h-1}$	$M/2$	$M/2 - 1$	$2(M/2)^{h-1}(M/2 - 1)$

$$\text{Total: } n \geq 1 + 2 \sum_{i=0}^{h-1} (M/2)^i (M/2 - 1) = 2(M/2)^h - 1$$

Therefore height of tree with  $n$  nodes is  $\Theta((\log n)/(\log M))$ .

# Analysis of B-tree operations

Assume each node stores its KVPs and child-pointers in a dictionary that supports  $O(\log M)$  search, insert, and delete.

Then *search*, *insert*, and *delete* work just like for 2-3 trees, and each require  $\Theta(\text{height})$  node operations.

Total cost is  $O\left(\frac{\log n}{\log M} \cdot (\log M)\right) = O(\log n)$ .

# Dictionaries in external memory

**Recall:** accessing a single location in *external memory* (e.g. hard disk) automatically loads a whole block (or “page”).

In an AVL tree or 2-3 tree,  $\Theta(\log n)$  pages are loaded in the worst case.

If  $M$  is small enough so an  $M$ -node fits into a single page, then a B-tree of order  $M$  only loads  $\Theta((\log n)/(\log M))$  pages.

This can result in a *huge* savings:  
memory access is often the largest time cost in a computation.

# B-tree variations

**Other strategies:** *insert* and *delete* without *backtracking* via *pre-emptive splitting* and *pre-emptive merging*.

**Red-black trees:** Identical to a B-tree with minsize 1 and maxsize 3, but each 2-node or 3-node is represented by 2 or 3 binary nodes, and each node holds a “color” value of red or black.

**B<sup>+</sup>-trees:** All KVPs are stored at the leaves (interior nodes just have keys), and the leaves are linked sequentially.



# Hashing in External Memory

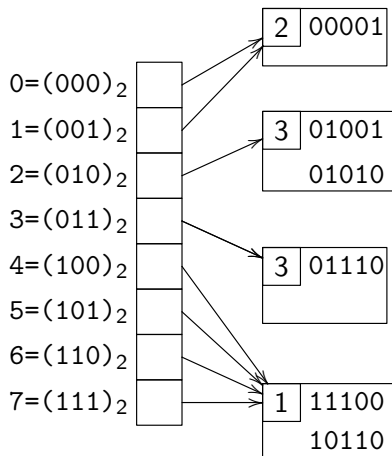
As before, if we have a *very large* dictionary that must be stored externally, how can we hash and minimize disk transfers?

Say external memory is stored in blocks (or “pages”) of size  $S$ . Most hash strategies access many pages (data is scattered).

Exception: **Linear Probing**. All hash table accesses will usually be in the same page. But  $\alpha$  must be kept small to avoid clustering, so there is a lot of wasted space.

New Idea: **Extendible Hashing**. Similar to a B-tree with height 1 and max size  $S$  at the leaves

# Extendible Hashing Overview



**Assumption:** Hash-function has values in  $\{0, 1, \dots, 2^L - 1\}$ .

The *directory* (similar to root node) is stored in *internal memory*.

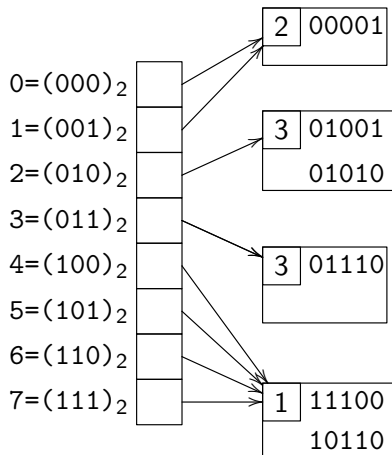
Contains array of size  $2^d$ , where  $d \leq L$  is called the *order*.

Each directory entry points to a *block* stored in *external memory*.

Each block contains at most  $S$  items. (Many entries can point to the same block.)

To look up a key  $k$  in the directory, use the  $d$  leading bits of  $h(k)$ .

# Extendible Hashing Details



Blocks are shared by entries in a specific manner:

- Every block  $B$  stores a *local depth*  $k_B \leq d$ .
- Hash values in  $B$  agree on leading  $k_B$  bits.
- All directory entries with the same  $k_B$  leading bits point to  $B$ .
- So  $2^{d-k_B}$  directory entries point to block  $B$ .

# Searching in extendible hashing

Searching is done in the directory, then in a block:

- Given a key  $k$ , compute  $h(k)$ .
- Leading  $d$  digits of  $h(k)$  give index in directory.
- Load block  $B$  at this index into main memory.
- Perform a search in  $B$  for all items with hash value  $h(k)$ .
- Search among them for the one with key  $k$ .

# Searching in extendible hashing

Searching is done in the directory, then in a block:

- Given a key  $k$ , compute  $h(k)$ .
- Leading  $d$  digits of  $h(k)$  give index in directory.
- Load block  $B$  at this index into main memory.
- Perform a search in  $B$  for all items with hash value  $h(k)$ .
- Search among them for the one with key  $k$ .

## Cost:

**CPU time:** depends on how the block are organized  
(hash table, balanced tree, sorted array)

**Disk transfers:** 1 (directory resides in internal memory)

# Insertion in Extensible Hashing

*insert*( $k, v$ ) is done as follows:

- Search for  $h(k)$  to find the proper block  $B$  for insertion
- If the  $B$  has space, then put  $(k, v)$  there.

# Insertion in Extensible Hashing

*insert*( $k, v$ ) is done as follows:

- Search for  $h(k)$  to find the proper block  $B$  for insertion
- If the  $B$  has space, then put  $(k, v)$  there.
- Else if the block is full and  $k_B < d$ , perform a *block split*:
  - ▶ Split  $B$  into two blocks  $B_0$  and  $B_1$ .
  - ▶ Separate items according to the  $(k_B + 1)$ -th bit.
  - ▶ Set local depth in  $B_0$  and  $B_1$  to  $k_B + 1$
  - ▶ Update references in the directory
  - ▶ Try again to insert

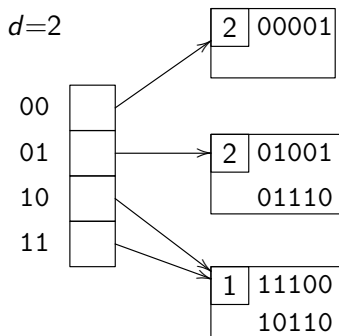
# Insertion in Extendible Hashing

*insert*( $k, v$ ) is done as follows:

- Search for  $h(k)$  to find the proper block  $B$  for insertion
- If the  $B$  has space, then put  $(k, v)$  there.
- Else if the block is full and  $k_B < d$ , perform a *block split*:
  - ▶ Split  $B$  into two blocks  $B_0$  and  $B_1$ .
  - ▶ Separate items according to the  $(k_B + 1)$ -th bit.
  - ▶ Set local depth in  $B_0$  and  $B_1$  to  $k_B + 1$
  - ▶ Update references in the directory
  - ▶ Try again to insert
- Else if the block is full and  $k_B = d$ , perform a *directory grow*:
  - ▶ Double the size of the directory ( $d \leftarrow d + 1$ )
  - ▶ Update references appropriately.
  - ▶ Then split block  $B$  (which is now possible).

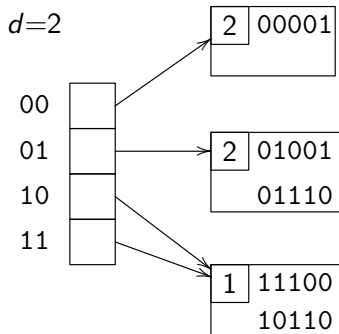


## Extendible hashing insert example with $S = 2$

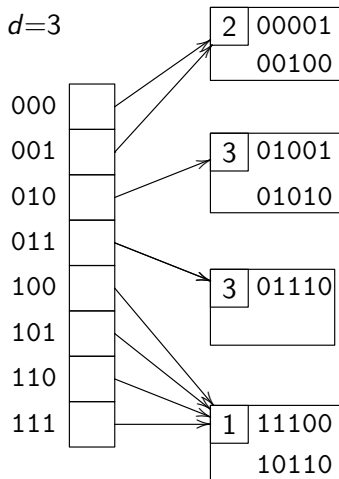


- Insert( 00100 )
- Insert( 01010 )

## Extendible hashing insert example with $S = 2$



- Insert( 00100 )
- Insert( 01010 )



## Extendible hashing conclusion

*delete*( $k$ ) is performed in a reverse manner to *insert*:

- Search for block  $B$  and remove  $k$  from it
- If block becomes too empty, then we perform a *block merge*
- If every block  $B$  has local depth  $k_B \leq d - 1$ , perform a *directory shrink*

But most likely just do *lazy deletion*.

# Extendible hashing conclusion

*delete*( $k$ ) is performed in a reverse manner to *insert*:

- Search for block  $B$  and remove  $k$  from it
- If block becomes too empty, then we perform a *block merge*
- If every block  $B$  has local depth  $k_B \leq d - 1$ , perform a *directory shrink*

But most likely just do *lazy deletion*.

Cost of *insert* and *delete*:

**CPU time:** Search in a block depends on the implementation

$\Theta(S)$  to do/undo one split

Directory grow/shrink costs  $\Theta(2^d)$  (but very rare).

**Disk transfers:** 1 when no split

# Summary of extendible hashing

- Directory is much smaller than total number of stored keys and should fit in main memory.
- To make more space, we only add one block.  
Rarely do we have to change the size of the directory.  
*Never* do we have to move all items in the dictionary (in contrast to normal hashing).
- Space usage is not too inefficient: can be shown that under uniform hashing, each block is expected to be 69% full.
- Potentially extra CPU cost